

---

# PartiQL User Guide

Ion Team

©2019 Amazon.com, Inc. or its affiliates.

# 1 Preface

The user's guide provides an explanation of the features implemented by PartiQL's kotlin implementation.

*This document is an early draft, contributions welcome!*

## 1.1 What is PartiQL

PartiQL is an implementation of [SQL++](#) based upon [lon's](#) type system. PartiQL is based on SQL92 and provides support for working with schemaless hierarchical data.

## 1.2 Conventions

*TBD*

## 1.3 Further Reading

- [SQL++](#)
- [lon's](#)

## 1.4 Bug Reports

We welcome you to use the GitHub issue tracker to report bugs or suggest features.

When filing an issue, please check existing open, or recently closed, issues to make sure somebody else hasn't already reported the issue. Please try to include as much information as you can. Details like these are incredibly useful:

- A reproducible test case or series of steps
- The version of our code being used
- Any modifications you've made relevant to the bug
- Anything unusual about your environment or deployment

## 1.5 Contribute

See our [contribute guide](#).

## 2 FAQ

TBD

## 3 Getting Started

PartiQL provides an interactive shell, or Read Evaluate Print Loop (REPL), that allows users to write and evaluate PartiQL queries.

### 3.1 Prerequisites

PartiQL requires the Java Runtime (JVM) to be installed on your machine. You can obtain the *latest* version of the Java Runtime from either

1. [OpenJDK](#), or [OpenJDK for Windows](#)
2. [Oracle](#)

Follow the instructions on how to set `JAVA_HOME` to the path where your Java Runtime is installed.

### 3.2 Download the PartiQL REPL

Each release of PartiQL comes with an archive that contains the PartiQL REPL as a zip file.

1. [Download](#). You may have to click on [Assets](#) to see the zip and tgz archives. the latest `partiql-cli`<sup>1</sup> zip archive to your machine.
2. Expand (unzip) the archive on your machine. Expanding the archive yields the following folder structure:

```
├─  
  partiql-cli  
    ├─ bin  
    │   ├─ partiql  
    │   └─ partiql.bat  
    └─ lib  
      └─ ...  
├─ README.md  
└─ Tutorial  
    └─ code
```

---

<sup>1</sup>The file will append PartiQL's release version to the archive, i.e., `partiql-cli-0.1.0.zip`.

```
|      └─ ...  
├─     tutorial.html  
└─     tutorial.pdf
```

where ... represents elided files/directories.

The root folder `partiql-cli` contains a `README.md` file and 3 subfolders

1. The folder `bin` contains startup scripts `partiql` for macOS and Unix systems and `partiql.bat` for Windows systems. Execute these files to start the REPL.
2. The folder `lib` contains all the necessary Java libraries needed to run PartiQL.
3. The folder `Tutorial` contains the tutorial in `pdf` and `html` form. The subfolder `code` contains 3 types of files:
  1. Data files with the extension `.env`. These files contains PartiQL data that we can query.
  2. PartiQL query files with the extension `.sql`. These files contain the PartiQL queries used in the tutorial.
  3. Sample query output files with the extension `.output`. These files contain sample output from running the tutorial queries on the appropriate data.

## 3.3 Running the PartiQL REPL

### 3.3.1 Windows

Run (double click on) `partiql.bat`. This should open a command-line prompt and start the PartiQL REPL which displays:

```
Welcome to the PartiQL REPL!  
PartiQL>
```

### 3.3.2 macOS (Mac) and Unix

1. Open a terminal and navigate to the `partiql-cli2` folder.
2. Start the REPL by typing `./bin/partiql` and pressing ENTER, which displays:

```
Welcome to the PartiQL REPL!  
PartiQL>
```

---

<sup>2</sup>The folder name will have the PartiQL version as a suffix, i.e., `partiql-cli-0.1.0`.

### 3.4 Testing the PartiQL REPL

Let's write a simple query to verify that our PartiQL REPL is working. At the `PartiQL>` prompt type:

```
PartiQL> SELECT * FROM [1,2,3]
```

and press `ENTER` twice. The output should look similar to:

```
<<
  {
    '_1': 1
  },
  {
    '_1': 2
  },
  {
    '_1': 3
  }
>>
```

Congratulations! You successfully installed and run the PartiQL REPL. The PartiQL REPL is now waiting for more input.

To exit the PartiQL REPL, press:

- `Control+D` in macOS or Unix
- `Control+C` on Windows

or close the terminal/command prompt window.

### 3.5 Loading data from a file

An easy way to load the necessary data into the REPL is use the `-e` switch when starting the REPL and provide the name of a file that contains your data.

```
./bin/partiql -e Tutorial/code/ql.env
```

You can then see what is loaded in the REPL's global environment using the **special** REPL command `!global_env`, i.e.,

```
PartiQL> !global_env;
```

```
{
  'hr': {
```

```
'employees': <<
  {
    'id': 3,
    'name': 'Bob Smith',
    'title': NULL
  },
  {
    'id': 4,
    'name': 'Susan Smith',
    'title': 'Dev Mgr'
  },
  {
    'id': 6,
    'name': 'Jane Smith',
    'title': 'Software Eng 2'
  }
>>
}
```

## 4 PartiQL CLI

### PartiQL CLI

Command line **interface for** executing PartiQL queries. Can be run in an interactive (REPL) mode or non-interactive.

#### Examples:

To run in REPL mode simply execute the executable without any arguments:

```
partiql
```

In non-interactive mode we use Ion as the format **for** input data which is bound to a global variable

named "input\_data", in the example below /logs/log.ion is bound to "input\_data":

```
partiql --query="SELECT * FROM input_data" --input=/logs/log.ion
```

The cli can output using PartiQL syntax or Ion using the --output-format option, e.g. to output binary ion:

```
partiql --query="SELECT * FROM input_data" --output-format=ION_BINARY --input=/logs/log.ion
```

To pipe input data in via stdin:

```
cat /logs/log.ion | partiql --query="SELECT * FROM input_data" --format=ION_BINARY >
output.10n
```

Option	Description
-----	-----
-e, --environment <File>	initial global environment (optional)
-h, --help	prints <b>this</b> help
-i, --input <File>	input file, requires the query option (optional)
-if, --input-format <InputFormat>	input format, requires the query option ( <b>default:</b> ION ) [ION, PARTIQL]
-w, --wrap-ion	wraps Ion input file values in a bag, requires the input format to be ION, requires the query option
-m, --monochrome	removes syntax highlighting <b>for</b> the REPL
-o, --output <File>	output file, requires the query option ( <b>default:</b> stdout)
-of, --output-format <OutputFormat>	output format, requires the query option ( <b>default:</b> PARTIQL) [PARTIQL, PARTIQL_PRETTY, ION_TEXT, ION_BINARY]
-p, --permissive	run the PartiQL query in PERMISSIVE typing mode
-q, --query <String>	PartiQL query, triggers non interactive mode

## 4.1 Building the CLI

The root Gradle build also builds the CLI. To build the CLI separately, execute:

```
./gradlew :cli:build
```

After building, distributable jars are located in the `cli/build/distributions` directory (relative to the project root).

Be sure to include the correct relative path to `gradlew` if you are not in the project root.

## 4.2 Using the CLI

The following command will build any dependencies before starting the CLI.

```
./gradlew :cli:run -q --args="<command line arguments>"
```

The CLI can be run in two manners, non-interactive and interactive (REPL).

## 4.3 REPL

To start an interactive read, eval, print loop (REPL) execute:

Note that running directly with Gradle will eat arrow keys and control sequences due to the Gradle daemon.

```
./cli/shell.sh
```

You will see a prompt that looks as follows:

```
Welcome to the PartiQL REPL!
PartiQL>
```

At this point, you can type in valid SQL/PartiQL and press enter *twice* to execute it:

```
PartiQL> SELECT id FROM `[{id: 5, name:"bill"}, {id: 6, name:"bob"}]` WHERE name = 'bob';
```

```
<<
{
  'id': 6
}
>>
```

Alternatively, you can denote the end of a query using a semi-colon:

```
PartiQL> SELECT id FROM `[{id: 5, name:"bill"}, {id: 6, name:"bob"}]` WHERE name = 'bob';
```

```
<<
{
  'id': 6
}
>>
```

The result of the previous expression is stored in the variable named `_`, so you can then run subsequent expressions based on the last one.

```
PartiQL> SELECT id + 4 AS name FROM _;
```

```
<<
{
  'name': 10
}
>>
```

Press control-D to exit the REPL.



### 4.3.1 Advanced REPL Features

To view the AST of a PartiQL statement, type the statement and press enter only *once*, then type `!!` and press enter:

```
PartiQL> 1 + 1
| !!
===
(
  ast
  (
    version
    1
  )
  (
    root
    (
      +
      (
        lit
        1
      )
      (
        lit
        1
      )
    )
  )
)
---
OK!
```

### 4.3.2 Initial Environment

The initial environment for the REPL can be setup with a configuration file, which should be a PartiQL file with a single `struct` containing the initial *global environment*.

For example, a file named `config.env` contains the following:

```
{
  'animals':[
    {'name': 'Kumo', 'type': 'dog'},
```

```
{'name': 'Mochi', 'type': 'dog'},
{'name': 'Lilikoi', 'type': 'unicorn'}
],
'types':[
  {'id': 'dog', 'is_magic': false},
  {'id': 'cat', 'is_magic': false},
  {'id': 'unicorn', 'is_magic': true}
]
}
```

The variables `animals` and `types` can both be bound to the execution environment for later access. To bind the environment file to the execution environment, start the REPL with the following command:

```
$ ./gradlew :cli:run -q --console=plain --args='-e config.env'
```

**Note:** Shell expansions such as `~` do not work within the value of the `args` argument.

Or, if you have extracted one of the compressed archives:

```
$ ./bin/partiql -e config.env
```

Expressions can then use the environment defined by `config.env`:

```
PartiQL> SELECT name, type, is_magic FROM animals, types WHERE type = id
```

```
<<
{
  'name': 'Kumo',
  'type': 'dog',
  'is_magic': false
},
{
  'name': 'Mochi',
  'type': 'dog',
  'is_magic': false
},
{
  'name': 'Lilikoi',
  'type': 'unicorn',
  'is_magic': true
}
>>
```

To see the current REPL environment you can use `!global_env`, for example for the file above:

```
PartiQL> !global_env;
```

```
{
  'types': [
    {
      'id': 'dog',
      'is_magic': false
    },
    {
      'id': 'cat',
      'is_magic': false
    },
    {
      'id': 'unicorn',
      'is_magic': true
    }
  ],
  'animals': [
    {
      'name': 'Kumo',
      'type': 'dog'
    },
    {
      'name': 'Mochi',
      'type': 'dog'
    },
    {
      'name': 'Lilikoi',
      'type': 'unicorn'
    }
  ]
}
```

You can also add new values to the global environment or replace existing values using `!add_to_global_env`. The example below replaces the value bound to `types`

```
PartiQL> !add_to_global_env {'types': []};
```

```
{
  'types': []
}
```

Let's look at what has changed:

```
PartiQL> !global_env
```

```
{
  'types': [],
  'animals': [
    {
      'name': 'Kumo',
      'type': 'dog'
    },
    {
      'name': 'Mochi',
      'type': 'dog'
    },
    {
      'name': 'Lilikoi',
      'type': 'unicorn'
    }
  ]
}
```

### 4.3.3 Working with Structure

Let's consider the following initial environment:

```
{
  'stores': [
    {
      'id': 5,
      'books': [
        {'title': 'A', 'price': 5.0, 'categories': ['sci-fi', 'action']},
        {'title': 'B', 'price': 2.0, 'categories': ['sci-fi', 'comedy']},
        {'title': 'C', 'price': 7.0, 'categories': ['action', 'suspense']},
        {'title': 'D', 'price': 9.0, 'categories': ['suspense']}
      ]
    },
    {
      'id': 6,
      'books': [
        {'title': 'A', 'price': 5.0, 'categories': ['sci-fi', 'action']},
        {'title': 'E', 'price': 9.5, 'categories': ['fantasy', 'comedy']}
      ]
    }
  ]
}
```

```
        {'title': 'F', 'price': 10.0, 'categories': ['history']}
    ]
}
]
```

Set the environment as below

```
PartiQL> !add_to_global_env { 'stores': [ { 'id': 5, 'books': [ {'title': 'A', 'price': 5.0,
    'categories': ['sci-fi', 'action']}, {'title': 'B', 'price': 2.0, 'categories': ['sci-fi',
    'comedy']}, {'title': 'C', 'price': 7.0, 'categories': ['action', 'suspense']}, {'title':
    'D', 'price': 9.0, 'categories': ['suspense']} ] }, { 'id': 6, 'books': [ {'title': 'A',
    'price': 5.0, 'categories': ['sci-fi', 'action']}, {'title': 'E', 'price': 9.5, '
    categories': ['fantasy', 'comedy']}, {'title': 'F', 'price': 10.0, 'categories': ['history
    ']} ] } ] }
```

If we wanted to find all books *as their own rows* with a price greater than 7 we can use paths on the `FROM` for this:

```
PartiQL> SELECT * FROM stores[*].books[*] AS b WHERE b.price > 7;
```

```
<<
{
  'title': 'D',
  'price': 9.0,
  'categories': [
    'suspense'
  ]
},
{
  'title': 'E',
  'price': 9.5,
  'categories': [
    'fantasy',
    'comedy'
  ]
},
{
  'title': 'F',
  'price': 10.0,
  'categories': [
    'history'
  ]
}
```

```
}  
>>
```

If you wanted to also de-normalize the store ID and title into the above rows:

```
PartiQL> SELECT s.id AS store, b.title AS title FROM stores AS s, @s.books AS b WHERE b.  
price > 7;
```

```
<<  
{  
  'store': 5,  
  'title': 'D'  
},  
{  
  'store': 6,  
  'title': 'E'  
},  
{  
  'store': 6,  
  'title': 'F'  
}  
>>
```

We can also use sub-queries with paths to predicate on sub-structure without changing the cardinality. So if we wanted to find all stores with books having prices greater than 9.5

```
PartiQL> SELECT * FROM stores AS s  
| WHERE EXISTS(  
|   SELECT * FROM @s.books AS b WHERE b.price > 9.5  
| );
```

```
<<  
{  
  'id': 6,  
  'books': [  
    {  
      'title': 'A',  
      'price': 5.0,  
      'categories': [  
        'sci-fi',  
        'action'  
      ]  
    }  
  ],  
},
```

```
{
  'title': 'E',
  'price': 9.5,
  'categories': [
    'fantasy',
    'comedy'
  ]
},
{
  'title': 'F',
  'price': 10.0,
  'categories': [
    'history'
  ]
}
]
}
>>
```

## 4.4 Reading/Writing Files

The REPL provides the `read_file` function to stream data from a file. The files need to be placed in the folder `cli`, and, if using the default file type (Ion), they must contain only a single Ion value (typically a list).

**Note:** Later on, we will introduce reading different file types, but we will first focus on the default (Ion).

For example, create a file called `data.ion` in the `cli` folder with the following contents

```
[
  { 'city': 'Seattle', 'state': 'WA' },
  { 'city': 'Bellevue', 'state': 'WA' },
  { 'city': 'Honolulu', 'state': 'HI' },
  { 'city': 'Rochester', 'state': 'NY' }
]
```

To select the cities that are in `HI` and `NY` states:

```
PartiQL> SELECT city FROM read_file('data.ion') AS c, `["HI", "NY"]` AS s WHERE c.state = s
;
```

```
<<
{
```

```
'city': 'Honolulu'
},
{
  'city': 'Rochester'
}
>>
```

The REPL also has the capability to write files with the `write_file` function:

```
PartiQL> write_file('out.ion', SELECT * FROM _);
```

A file called `out.ion` will be created in the `cli` directory with the following contents:

```
$partiql_bag::[
  {
    city: Honolulu
  },
  {
    city: Rochester
  }
]
```

Notice that PartiQL added the annotation of `$partiql_bag` to the Ion list. When outputting to Ion, we use type annotations to represent some PartiQL values/types not in Ion.

Functions and expressions can be used in the *global configuration* as well. Consider the following `config.ion`:

```
{
  'data': read_file('data.ion')
}
```

The `data` variable will now be bound to file containing Ion:

```
PartiQL> SELECT * FROM data;
```

```
<<
{
  'city': 'Seattle',
  'state': 'WA'
},
{
  'city': 'Bellevue',
  'state': 'WA'
}
```



```
    },  
    {  
      'city': 'Honolulu',  
      'state': 'HI'  
    },  
    {  
      'city': 'Rochester',  
      'state': 'NY'  
    }  
  ]  
>>
```

## 5 TSV/CSV Data

The `read_file` function supports an optional `struct` argument to add additional parsing options. Parsing delimited files can be specified with the `type` field with a string `tsv` or `csv` to parse tab or comma separated values respectively.

**Note:** One might ask how this differs from reading in Ion files. With Ion files, PartiQL expects a *single* value – typically a bag/list, but it can also be literals such as strings or integers. CSV & TSV rows, on the other hand, are *always* interpreted as being contained within a bag.

Create a file called `simple.csv` in the `cli` directory with the following contents:

```
title,category,price  
harry potter,book,7.99  
dot,electronics,49.99  
echo,electronics,99.99
```

You can read the file with the following CLI command:

```
PartiQL> read_file('simple.csv', {'type':'csv'});
```

```
<<  
  {  
    _0: 'title',  
    _1: 'category',  
    _2: 'price'  
  },  
  {  
    _0: 'harry potter',  
    _1: 'book',  
    _2: '7.99'
```

```
    },
    {
      _0: 'dot',
      _1: 'electronics',
      _2: '49.99'
    },
    {
      _0: 'echo',
      _1: 'electronics',
      _2: '99.99'
    }
  ]
}>>
```

The options `struct` can also define if the first row for delimited data should be the column names with the `header` field.

```
PartiQL> read_file('simple.csv', {'type': 'csv', 'header': true});
```

```
<<
{
  'title': 'harry potter',
  'category': 'book',
  'price': '7.99'
},
{
  'title': 'dot',
  'category': 'electronics',
  'price': '49.99'
},
{
  'title': 'echo',
  'category': 'electronics',
  'price': '99.99'
}
}>>
```

Auto-conversion for numeric and timestamp values can also be specified as follows:

```
PartiQL> read_file('simple.csv', {'type': 'csv', 'header': true, 'conversion': 'auto'});
```

```
<<
{
  'title': 'harry potter',
```

```
'category': 'book',
'price': 7.99
},
{
  'title': 'dot',
  'category': 'electronics',
  'price': 49.99
},
{
  'title': 'echo',
  'category': 'electronics',
  'price': 99.99
}
>>
```

Writing TSV/CSV data can be done by specifying the optional `struct` argument to specify output format to the `write_file` function. Similar to the `read_file` function, the `type` field can be used to specify `tsv`, `csv`, or `ion` output.

```
PartiQL> write_file('out.tsv', SELECT name, type FROM animals, {'type':'tsv'});
```

This would produce the following file:

```
$ cat out.tsv
Kumo    dog
Mochi   dog
Lilikoi unicorn
```

The options `struct` can also specify a `header` Boolean field to indicate whether the output TSV/CSV should have a header row.

```
PartiQL> write_file('out.tsv', SELECT name, type FROM animals, {'type':'tsv', 'header':true
  });
```

Which would produce the following file:

```
$ cat out.tsv
name    type
Kumo    dog
Mochi   dog
Lilikoi unicorn
```

## 5.1 Predefined CSV Data

The `read_file` function provides options to read other predefined CSV data formats. For example, if a CSV file is exported from PostgreSQL, we can use the following command to read the file:

```
PartiQL> read_file('simple_postgresql.csv', {'type': 'postgresql_csv'})
```

Other available options for the argument `type` besides `postgresql_csv` are `excel_csv`, `mysql_csv`, and `postgresql_text`.

## 5.2 Customized CSV Data

The `read_file` function also provides options to read customized CSV data formats. For example, we have a data file where the whitespace is the separator as shown below:

```
title category price
harry_potter book 7.99
dot electronics 49.99
echo electronics 99.99
```

We can use the following command to read the file:

```
PartiQL> read_file('customized.csv', {'type': 'customized', 'delimiter': ' ', 'header': true})
```

All the available options for customized CSV files are as follows: 1. Ignore empty lines: `'ignore_empty_line': true` 2. Ignore spaces surrounding comma: `'ignore_surrounding_space': true` 3. Trim leading and trailing blanks: `'trim': true` 4. Set line breaker (only working with `'\r'`, `'\n'` and `'\r\n'`): `'line_breaker': '\n'` 5. Set escape sign (single character only): `'escape': '\'` 6. Set quote sign (single character only): `'quote': '\"'` 7. Set delimiter sign (single character only): `'delimiter': ','`

## 5.3 Querying Amazon DynamoDB

We also provide a CLI function, `query_ddb`, that allows you to query AWS DynamoDB tables and perform additional computations on the response.

**Note:** This implementation uses the [Default Credentials Provider](#). Please see the link to determine how you can specify which account/profile to use.

For example, let's say you have a DDB table called `users` with primary-key of `id`. If your AWS credentials and configurations are set correctly, you can perform the following:

```
PartiQL> query_ddb('SELECT * FROM users WHERE id = 0');
```

Which, depending on the content of your table, will return something like:

```
[
  {
    'id': 0,
    'name': 'John Doe',
    'age': 22
  }
]
```

Now, while we don't recommend performing scans on your DDB tables (due to the cost), you can execute queries similar to:

```
PartiQL> !add_to_global_env {
          'fullNames': <<
            'John Doe',
            'Sarah Jane',
            'Boe Jackson'
          >>
        };
PartiQL> SELECT VALUE id
        FROM query_ddb('SELECT id, name, age FROM users')
        WHERE name IN fullNames;
```

The above query will get the `id`'s of all the users in your local environment, something like:

```
<<
  0,
  17,
  1004
>>
```

Also, if you'd like to perform insertions into DDB tables, you can perform them as follows:

```
PartiQL> query_ddb('INSERT INTO users VALUE {'id': 96, 'name': 'Kim Lu', 'age': 26}
                  ');
```

**Note:** You can escape the PartiQL single-quote by prepending another single-quote. See above.

For in-depth documentation on valid DDB PartiQL queries, please reference the official [AWS DynamoDB PartiQL Docs](#).

## 5.4 Permissive Typing Mode

By default, the CLI/REPL runs in [LEGACY](#) typing mode, which will give an evaluation time error in the case of data type mismatches.

```
# Running in the default LEGACY typing mode
PartiQL> 1 + 'foo';
org.partiql.lang.eval.EvaluationException: ...
...
```

Specifying the `-p` or `-permissive` flag will allow you to run PartiQL queries in [PERMISSIVE](#) typing mode, which will return [MISSING](#) in the case of data type mismatches.

```
# Running in PERMISSIVE typing mode
PartiQL> 1 + 'foo';
===
MISSING
---
OK!
```

## 6 PartiQL Tutorial

TBD

## 7 PartiQL User Guide

### 7.1 Introduction

This is the PartiQL implementation's user guide. The goal of this document is to provide to users of PartiQL information on the features implemented and any deviation from the PartiQL specification.

### 7.2 Data Types

#### 7.2.1 Decimal

PartiQL decimals are based on [lon decimals](#) from the [lon Specification](#)<sup>[1]</sup> but with a maximum precision of 38 digits, numbers outside this precision range will be rounded using a round [half even strategy](#). Examples:



- `null`
- `integer`
- `boolean`
- `float`
- `decimal`
- `timestamp`
- `symbol`
- `string`
- `list`
- `struct`
- `bag`

**Header** `CAST(exp AS dt)`

**Purpose** Given an expression, `exp` and the data type name, `dt`, evaluate `expr` to a value, `v` and alter the data type of `v` to `DT(dt)`. If the conversion cannot be made the implementation signals an error.

The runtime support for casts is

- Casting to `null` from
  - `null` is a no-op
  - `missing` returns `null`
  - else error
- Casting to `missing` from
  - `missing` is a no-op
  - `null` returns `missing`
  - else error
- Casting to `integer` from
  - Integer: is a no-op
  - Boolean: `true` returns 1, `false` returns 0
  - String or Symbol: attempt to parse the content as an Integer and return the Integer, else error.
  - Float or Decimal: gets narrowed to Integer
  - else error
- Casting to `boolean` from
  - Boolean is a no-op
  - Integer or Decimal or Float: if `v` is a representation of the number 0 (e.g., 0 or -0 or 0e0 or 0d0 et.) then `false` else `true`
  - String or Symbol: `true` unless `v` matches-ignoring character case-the lon string "`false`" or matches-ignoring character case- the lon symbol '`false`' then return `false`
  - else error
- Casting to `float` from



- Float is a no-op
- Boolean: **false** return 0.0, **true** returns 1.0
- Integer or Decimal: convert to Float and return
- String or Symbol: attempt to parse as Float and return the Float value, else error
- else error
- Casting to **decimal** from
  - Decimal is a no-op
  - Boolean: return 1d0 if **true**, 0d0 if **false**
  - String or Symbol: attempt to parse as Decimal and return Decimal value, else error
  - else error
- Casting to **timestamp** from
  - Timestamp is a no-op
  - String or Symbol: attempt to parse as Timestamp and return the Timestamp value, else error
  - else error
- Casting to **symbol** from
  - Symbol is a no-op
  - Integer or Float or Decimal: narrow to Integer and return the value as a Symbol, i.e., a Symbol with the same sequence of digits as characters
  - String: return the String as a Symbol, i.e., represent the same sequence of characters as a Symbol
  - Boolean: return '**true**' for **true** and '**false**' for **false**
  - Timestamp: return the Symbol representation of the Timestamp, i.e., represent the same sequence of digits and characters as a Symbol
  - else error
- Casting to **string** from
  - String is a no-op
  - Integer or Float or Decimal: narrow to Integer and return the value as a String, i.e., a String with the same sequence of digits as characters
  - Symbol: return the String as a Symbol, i.e., represent the same sequence of characters as a String
  - Boolean: return "**true**" for **true** and "**false**" for **false**
  - Timestamp: return the String representation of the Timestamp, i.e., represent the same sequence of digits and characters as a String
  - else error
- Casting to **list** from
  - List is a no-op
  - Bag: return a list with the same elements. The order of the elements in the resulting list is unspecified.

- else error
- Casting to `struct` from
  - Struct is a no-op
  - else error
- Casting to `bag` from
  - Bag is a no-op
  - List: return a list with the same elements. The order of the elements in the resulting bag is unspecified.

Examples :

```
-- Unknowns propagation
CAST(null AS null) -- null
CAST(missing AS null) -- null
CAST(missing AS missing) -- null
CAST(null AS missing) -- null
CAST(null AS boolean) -- null (null AS any data type name result to null)
CAST(missing AS boolean) -- null (missing AS any data type name result to null)

-- any value that is not an unknown cannot be cast to `null` or `missing`
CAST(true AS null) -- error
CAST(true AS missing) -- error
CAST(1 AS null) -- error
CAST(1 AS missing) -- error

-- AS boolean
CAST(true AS boolean) -- true no-op
CAST(0 AS boolean) -- false
CAST(1 AS boolean) -- true
CAST(`1e0` AS boolean) -- true (float)
CAST(`1d0` AS boolean) -- true (decimal)
CAST(`a` AS boolean) -- false
CAST(`true` AS boolean) -- true (PartiQL string 'true')
CAST(`true` AS boolean) -- true (Ion symbol `true`)
CAST(`false` AS boolean) -- false (Ion symbol `false`)

-- AS integer
CAST(true AS integer) -- 1
CAST(false AS integer) -- 0
CAST(1 AS integer) -- 1
CAST(`1d0` AS integer) -- 1
CAST(`1d3` AS integer) -- 1000
CAST(1.00 AS integer) -- 1
```

```

CAST('12' AS integer) -- 12
CAST('aa' AS integer) -- error
CAST(`'22'` AS integer) -- 22
CAST(`'x'` AS integer) -- error

-- AS flaot
CAST(true AS float) -- 1e0
CAST(false AS float) -- 0e0
CAST(1 AS float) -- 1e0
CAST(`1d0` AS float) -- 1e0
CAST(`1d3` AS float) -- 1000e0
CAST(1.00 AS float) -- 1e0
CAST('12' AS float) -- 12e0
CAST('aa' AS float) -- error
CAST(`'22'` AS float) -- 22e0
CAST(`'x'` AS float) -- error

-- AS decimal
CAST(true AS decimal) -- 1.
CAST(false AS decimal) -- 0.
CAST(1 AS decimal) -- 1.
CAST(`1d0` AS decimal) -- 1. (REPL printer serialized to 1.)
CAST(`1d3` AS decimal) -- 1d3
CAST(1.00 AS decimal) -- 1.00
CAST('12' AS decimal) -- 12.
CAST('aa' AS decimal) -- error
CAST(`'22'` AS decimal) -- 22.
CAST(`'x'` AS decimal) -- error

-- AS timestamp
CAST(`2001T` AS timestamp) -- 2001T
CAST('2001-01-01T' AS timestamp) -- 2001-01-01
CAST(`'2010-01-01T00:00:00.000Z'` AS timestamp) -- 2010-01-01T00:00:00.000Z
CAST(true AS timestamp) -- error
CAST(2001 AS timestamp) -- error

-- AS symbol
CAST(`'xx'` AS symbol) -- xx (`'xx'` is an Ion symbol)
CAST('xx' AS symbol) -- xx ('xx' is a string)
CAST(42 AS symbol) -- '42'
CAST(`1e0` AS symbol) -- '1'
CAST(`1d0` AS symbol) -- '1'
CAST(true AS symbol) -- 'true'

```

```

CAST(false                AS symbol) -- 'false'
CAST(`2001T`              AS symbol) -- '2001T'
CAST(`2001-01-01T00:00:00.000Z` AS symbol) -- '2001-01-01T00:00:00.000Z`

-- AS string
CAST(`'xx'`               AS string) -- "xx" (`'xx'` is an Ion symbol)
CAST('xx'                 AS string) -- "xx" ('xx' is a string)
CAST(42                    AS string) -- "42"
CAST(`1e0`                 AS string) -- "1.0"
CAST(`1d0`                 AS string) -- "1"
CAST(true                  AS string) -- "true"
CAST(false                 AS string) -- "false"
CAST(`2001T`              AS string) -- "2001T"
CAST(`2001-01-01T00:00:00.000Z` AS string) -- "2001-01-01T00:00:00.000Z"

-- AS struct
CAST(`{ a: 1 }` AS struct) -- { a:1 }
CAST(true      AS struct) -- err

-- AS list
CAST(`[1, 2, 3]` AS list) -- [ 1, 2, 3 ] (REPL does not display the parens and commas)
CAST(<<'a', { 'b':2 }>> AS list) -- [ a, { 'b':2 } ] (REPL does not display the parens and commas)
CAST({ 'b':2 } AS list) -- error

-- AS bag
CAST([1,2,3] AS bag) -- <<1,2,3>> (REPL does not display << >> and commas)
CAST([1,[2],3] AS bag) -- <<1,[2],3>> (REPL does not display << >> and commas)
CAST(<<'a', 'b'>> AS bag) -- <<'a', 'b'>> (REPL does not display << >> and commas)

```

### 7.3.3 CHAR\_LENGTH, CHARACTER\_LENGTH

Counts the number of characters in the specified string, where ‘character’ is defined as a single unicode code point.

*Note:* CHAR\_LENGTH and CHARACTER\_LENGTH are synonyms.

**Signature** CHAR\_LENGTH: String -> Integer

CHARACTER\_LENGTH: String -> Integer

**Header** CHAR\_LENGTH(str)

`CHARACTER_LENGTH(str)`

**Purpose** Given a `String` value `str` return the number of characters (code points) in `str`.

**Examples**

```
CHAR_LENGTH('')           -- 0
CHAR_LENGTH('abcdefg')    -- 7
CHAR_LENGTH('😄😄😄😄') -- 4 (non-BMP unicode characters)
CHAR_LENGTH('🇬🇧')         -- 2 (because '🇬🇧' is two codepoints: the letter 'e' and
                           combining character U+032B)
```

### 7.3.4 COALESCE

Evaluates the arguments in order and returns the first non unknown, i.e. first non-`null` or non-`missing`. This function does **not** propagate `null` and `missing`.

**Signature** `COALESCE: Any Any ... -> Any`

**Header** `COALESCE(exp, [exp ...])`

**Purpose** Given a list of 1 or more arguments, evaluates the arguments left-to-right and returns the first value that is **not** an unknown (`missing` or `null`).

**Examples**

```
COALESCE(1)               -- 1
COALESCE(null)             -- null
COALESCE(null, null)       -- null
COALESCE(missing)          -- null
COALESCE(missing, missing) -- null
COALESCE(1, null)          -- 1
COALESCE(null, null, 1)    -- 1
COALESCE(null, 'string')   -- 'string'
COALESCE(missing, 1)       -- 1
```

### 7.3.5 DATE\_ADD

Given a data part, a quantity and a timestamp, returns an updated timestamp by altering datetime part by quantity

**Signature** `DATE_ADD: DateTimePart Integer Timestamp -> Timestamp`

Where `DateTimePart` is one of

- `year`

- `month`
- `day`
- `hour`
- `minute`
- `second`

**Header** `DATE_ADD(dp, q, timestamp)`

**Purpose** Given a data part `dp`, a quantity `q`, and, an `Ion` timestamp `timestamp` returns an updated timestamp by applying the value for `q` to the `dp` component of `timestamp`. Positive values for `q` add to the `timestamp`'s `dp`, negative values subtract.

The value for `timestamp` as well as the return value from `DATE_ADD` must be a valid [Ion Timestamp](#)

### Examples

```
DATE_ADD(year, 5, `2010-01-01T`)           -- 2015-01-01 (equivalent to 2015-01-01T)
DATE_ADD(month, 1, `2010T`)                 -- 2010-02T (result will add precision as
necessary)
DATE_ADD(month, 13, `2010T`)                 -- 2011-02T
DATE_ADD(day, -1, `2017-01-10T`)             -- 2017-01-09 (equivalent to 2017-01-09T)
DATE_ADD(hour, 1, `2017T`)                   -- 2017-01-01T01:00:00:00
DATE_ADD(hour, 1, `2017-01-02T03:04Z`)       -- 2017-01-02T04:04Z
DATE_ADD(minute, 1, `2017-01-02T03:04:05.006Z`) -- 2017-01-02T03:05:05.006Z
DATE_ADD(second, 1, `2017-01-02T03:04:05.006Z`) -- 2017-01-02T03:04:06.006Z
```

## 7.3.6 DATE\_DIFF

Given a datetime part and two valid timestamps returns the difference in datetime parts.

**Signature** `DATE_DIFF: DateTimePart Timestamp Timestamp -> Integer`

See [DATE\\_ADD](#) for the definition of `DateTimePart`

**Header** `DATE_DIFF(dp, t1, t2)`

**Purpose** Given a datetime part `dp` and two timestamps `t1` and `t2` returns the difference in value for `dp` part of `t1` with `t2`. The return value is a negative integer when the `dp` value of `t1` is greater than the `dp` value of `t2`, and, a positive integer when the `dp` value of `t1` is less than the `dp` value of `t2`.

### Examples

```
DATE_DIFF(year, `2010-01-01T`, `2011-01-01T`) -- 1
DATE_DIFF(year, `2010T`, `2010-05T`)           -- 4 (2010T is equivalent to
2010-01-01T00:00:00.000Z)
DATE_DIFF(month, `2010T`, `2011T`)              -- 12
DATE_DIFF(month, `2011T`, `2010T`)              -- -12
```

```
DATE_DIFF(day, `2010-01-01T23:00T`, `2010-01-02T01:00T`) -- 0 (need to be at least 24h
apart to be 1 day apart)
```

### 7.3.7 EXISTS

Given a PartiQL value returns **true** if and only if the value is a non-empty container(bag, sexp, list or struct), returns **false** otherwise.

**Signature** `EXISTS: Container -> Boolean`

**Header** `EXISTS(val)`

**Purpose** Given a PartiQL value `val`, if `val` is

1. a container with size > 0, `EXISTS(val)` returns **true**;
2. a container with size = 0, `EXISTS(val)` returns **false**;
3. not a container, `EXISTS(val)` throws an error.

Note: 1. This function does **not** propagate **null** and **missing**.

#### Examples

```
EXISTS(`[]`)      -- false (empty list)
EXISTS(`[1, 2, 3]`) -- true (non-empty list)
EXISTS(`[missing]`) -- true (non-empty list)
EXISTS(`{}`)      -- false (empty struct)
EXISTS(`{ a: 1 }`) -- true (non-empty struct)
EXISTS(`()`)      -- false (empty s-expression)
EXISTS(`(+ 1 2)`) -- true (non-empty s-expression)
EXISTS(`<<>>`)    -- false (empty bag)
EXISTS(`<<null>>`) -- true (non-empty bag)
EXISTS(1)         -- error
EXISTS(`2017T`)   -- error
EXISTS(null)      -- error
EXISTS(missing)   -- error
```

### 7.3.8 EXTRACT

Given a datetime part and a datetime type returns then datetime's datetime part value.

**Signature** `EXTRACT: ExtractDateTimePart DateTime -> Integer`

where `ExtractDateTimePart` is one of

- `year`
- `month`
- `day`

- hour
- minute
- second
- timezone\_hour
- timezone\_minute

and `DateTime` type is one of

- DATE
- TIME
- TIMESTAMP

Note that `ExtractDateTimePart` **differs** from `DateTimePart` in `DATE_ADD`.

**Header** `EXTRACT(edp FROM t)`

**Purpose** Given a datetime part, `edp`, and a datetime type `t` return `t`'s value for `edp`. This function allows for `t` to be unknown (`null` or `missing`) but **not** `edp`. If `t` is unknown the function returns `null`.

### Examples

```
EXTRACT(YEAR FROM `2010-01-01T`)           -- 2010
EXTRACT(MONTH FROM `2010T`)                 -- 1 (equivalent to 2010-01-01
      T00:00:00.000Z)
EXTRACT(MONTH FROM `2010-10T`)              -- 10
EXTRACT(HOUR FROM `2017-01-02T03:04:05+07:08`) -- 3
EXTRACT(MINUTE FROM `2017-01-02T03:04:05+07:08`) -- 4
EXTRACT(TIMEZONE_HOUR FROM `2017-01-02T03:04:05+07:08`) -- 7
EXTRACT(TIMEZONE_MINUTE FROM `2017-01-02T03:04:05+07:08`) -- 8
EXTRACT(YEAR FROM DATE '2010-01-01')        -- 2010
EXTRACT(MONTH FROM DATE '2010-01-01')       -- 1
EXTRACT(DAY FROM DATE '2010-01-01')         -- 1
EXTRACT(HOUR FROM DATE '2010-01-01')        -- 0
EXTRACT(MINUTE FROM DATE '2010-01-01')      -- 0
EXTRACT(SECOND FROM DATE '2010-01-01')      -- 0
EXTRACT(HOUR FROM TIME '23:12:59')           -- 23
EXTRACT(MINUTE FROM TIME '23:12:59')        -- 12
EXTRACT(SECOND FROM TIME '23:12:59')        -- 59
EXTRACT(SECOND FROM TIME (2) '23:12:59.128') -- 59.13
EXTRACT(TIMEZONE_HOUR FROM TIME WITH TIME ZONE '23:12:59-08:30') -- -8
EXTRACT(TIMEZONE_MINUTE FROM TIME WITH TIME ZONE '23:12:59-08:30') -- -30
```

Note that `timezone_hour` and `timezone_minute` are **not supported** for `DATE` and `TIME` (without time zone) type.



### 7.3.9 FILTER\_DISTINCT

**Signature** `FILTER_DISTINCT: Container -> Bag`

**Header** `FILTER_DISTINCT(c)`

**Purpose** Returns a bag of distinct values contained within a bag, list, sexp, or struct. If the container is a struct, the field names are not considered.

Examples :

```
FILTER_DISTINCT([0, 0, 1])           -- <<0, 1>>
FILTER_DISTINCT(<<0, 0, 1>>)         -- <<0, 1>>
FILTER_DISTINCT(SEXP(0, 0, 1))       -- <<0, 1>>
FILTER_DISTINCT({'a': 0, 'b': 0, 'c': 1}) -- <<0, 1>>
```

### 7.3.10 LOWER

Given a string convert all upper case characters to lower case characters.

**Signature** `LOWER: String -> String`

**Header** `LOWER(s)`

**Purpose** Given a string, *s*, alter every upper case character in *s* to lower case. Any non-upper cased characters remain unchanged. This operation does rely on the locale specified by the runtime configuration. The implementation, currently, relies on Java's [String.toLowerCase\(\)](#) documentation.

**Examples**

```
LOWER('AbCdEfG!@#') -- 'abcdefg!@#'
```

### 7.3.11 MAKE\_DATE

Given the integer values for the year, month and day returns the associated date.

**Signature** `MAKE_DATE: Year_Int Month_Int Day_Int -> Date`

where *Year\_Int*, *Month\_Int*, *Day\_Int* are the Integers representing *year*, *month* and *day* respectively for the date.

**Header** `MAKE_DATE(year, month, day)`

**Purpose** Given integer values for *year*, *month* and *day*, returns an associated *Date*. This function allows arguments to be unknowns i.e. (*null* or *missing*).

Examples :

```
MAKE_DATE(2021, 02, 28)           -- 2021-02-28
MAKE_DATE(2020, 02, 29)           -- 2020-02-29
MAKE_DATE(2021, 12, 31)           -- 2021-12-31
MAKE_DATE(null, 02, 28)           -- null
MAKE_DATE(2021, null, 28)          -- null
MAKE_DATE(2021, 02, null)          -- null
MAKE_DATE(missing, 02, 28)         -- null
MAKE_DATE(2021, missing, 28)       -- null
MAKE_DATE(2021, 02, missing)       -- null
```

### 7.3.12 MAKE\_TIME

Given the values for the hour (int), minute (int), second (BigDecimal) and optionally for timezone minutes (int) returns the associated time.

**Signature** `MAKE_TIME: Hour_Int Minute_Int Second_BigDecimal (optional)TimezoneMinutes_Int -> Time`

where `Hour_Int`, `Minute_Int`, `TimezoneMinutes_Int` are the Integers representing `hour`, `minute` and `timezone minutes` respectively for the time. `Second_BigDecimal` is the BigDecimal representing the `second` with fraction of the time.

**Header** `MAKE_TIME(hour, minute, second, timezoneMinutes?)`

**Purpose** Given values for `hour` (int), `minute` (int), `second` (BigDecimal) and optionally for `timezone minutes` (int), returns an associated `Time`. This function allows arguments to be unknowns i.e. (`null` or `missing`).

Examples :

```
MAKE_TIME(23, 59, 59.)           -- 23:59:59
MAKE_TIME(23, 59, 59.12345, 330) -- 23:59:59.12345+05:30
MAKE_TIME(23, 59, 59.12345, -330) -- 23:59:59.12345-05:30
MAKE_TIME(null, 02, 28.)          -- null
MAKE_TIME(12, null, 28.)           -- null
MAKE_TIME(21, 02, null)            -- null
MAKE_TIME(missing, 02, 28.)        -- null
MAKE_TIME(23, missing, 59.)        -- null
MAKE_TIME(21, 02, missing)         -- null
MAKE_TIME(21, 02, 28., null)       -- null
MAKE_TIME(21, 02, 28., missing)    -- null
```

### 7.3.13 SIZE

Given any container data type (i.e., list, structure or bag) return the number of elements in the container.

**Signature** `SIZE: Container -> Integer`

**Header** `SIZE(c)`

**Purpose** Given a container, `c`, return the number of elements in the container. If the input to `SIZE` is not a container the implementation throws an error.

#### Examples

```
SIZE(`[]`)           -- 0
SIZE(`[null]`)       -- 1
SIZE(`[1,2,3]`)      -- 3
SIZE(<<'foo', 'bar'>>) -- 2
SIZE(`{foo: bar}`)   -- 1 (number of key-value pairs)
SIZE(`[{foo: 1}, {foo: 2}]`) -- 2
SIZE(12)             -- error
```

### 7.3.14 NULLIF

Given two expressions return `null` if the two expressions evaluate to the same value, else, returns the result of evaluating the first expression

**Signature** `NULLIF: Any Any -> Any`

**Header** `NULLIF(e1, e2)`

**Purpose** Given two expression, `e1` and `e2`, evaluate both expression to get `v1` and `v2` respectively, and return `null` if and only if `v1` equals `v2`, else, return `v1`. The implementation of `NULLIF` uses `=` for equality, i.e., `v1` and `v2` are considered equal by `NULLIF` if and only if `v1 = v2` is true.

*Note*, `NULLIF` does **not** propagate unknowns (`null` and `missing`).

#### Examples

```
NULLIF(1, 1)         -- null
NULLIF(1, 2)         -- 1
NULLIF(1.0, 1)       -- null
NULLIF(1, '1')       -- 1
NULLIF([1], [1])     -- null
NULLIF(1, NULL)      -- 1
NULLIF(NULL, 1)      -- null
NULLIF(null, null)   -- null
NULLIF(missing, null) -- null
NULLIF(missing, missing) -- null
```

### 7.3.15 SUBSTRING

Given a string, a start index and optionally a length, returns the substring from the start index up to the end of the string, or, up to the length provided.

**Signature** `SUBSTRING: String Integer [ NNegInteger ] -> String`

Where `NNegInteger` is a non-negative integer, i.e., 0 or greater.

**Header** `SUBSTRING(str, start [ , length ])`

`SUBSTRING(str FROM start [ FOR length ])`

**Purpose** Given a string, `str`, a start position, `start` and optionally a length, `length`, extract the characters (code points) starting at index `start` and ending at  $(start + length) - 1$ . If `length` is omitted, then proceed till the end of `str`.

The first character of `str` has index 1.

#### Examples

```
SUBSTRING("123456789", 0)      -- "123456789"
SUBSTRING("123456789", 1)      -- "123456789"
SUBSTRING("123456789", 2)      -- "23456789"
SUBSTRING("123456789", -4)     -- "123456789"
SUBSTRING("123456789", 0, 999) -- "123456789"
SUBSTRING("123456789", 0, 2)   -- "1"
SUBSTRING("123456789", 1, 999) -- "123456789"
SUBSTRING("123456789", 1, 2)   -- "12"
SUBSTRING("1", 1, 0)           -- ""
SUBSTRING("1", 1, 0)           -- ""
SUBSTRING("1", -4, 0)          -- ""
SUBSTRING("1234", 10, 10)      -- ""
```

### 7.3.16 TO\_STRING

Given a timestamp and a format pattern return a string representation of the timestamp in the given format.

**Signature** `TO_STRING: Timestamp TimeFormatPattern -> String`

Where `TimeFormatPattern` is a String with the following special character interpretations

---

Format	Example	Description
<code>yy</code>	69	2-digit year

---

Format	Example	Description
y	1969	4-digit year
yyyy	1969	Zero padded 4-digit year
M	1	Month of year
MM	01	Zero padded month of year
MMM	Jan	Abbreviated month year name
MMMM	January	Full month of year name
MMMMM	J	Month of year first letter (NOTE: not valid for use with to_timestamp function)
d	2	Day of month (1-31)
dd	02	Zero padded day of month (01-31)
a	AM	AM or PM of day
h	3	Hour of day (1-12)
hh	03	Zero padded hour of day (01-12)
H	3	Hour of day (0-23)
HH	03	Zero padded hour of day (00-23)
m	4	Minute of hour (0-59)
mm	04	Zero padded minute of hour (00-59)
s	5	Second of minute (0-59)
ss	05	Zero padded second of minute (00-59)
S	0	Fraction of second (precision: 0.1, range: 0.0-0.9)
SS	06	Fraction of second (precision: 0.01, range: 0.0-0.99)
SSS	060	Fraction of second (precision: 0.001, range: 0.0-0.999)
...	...	...
SSSSSSSS	060000000	Fraction of second (maximum precision: 1 nanosecond, range: 0.0-0.999999999)
n	600000000	Nano of second
X	+07 or Z	Offset in hours or "Z" if the offset is 0
XX or XXXX	+0700 or Z	Offset in hours and minutes or "Z" if the offset is 0

Format	Example	Description
XXX or XXXXX	+07:00 or Z	Offset in hours and minutes or “Z” if the offset is 0
x	+07	Offset in hours
xx or xxxx	+0700	Offset in hours and minutes
xxx or xxxxx	+07:00	Offset in hours and minutes

**Header** `TO_STRING(t, f)`

**Purpose** Given a timestamp, `t`, and a format pattern, `f`, as a String, return a string representation of `t` in format `f`.

**Examples**

```

TO_STRING(`1969-07-20T20:18Z`, 'MMMM d, y')           -- "July 20, 1969"
TO_STRING(`1969-07-20T20:18Z`, 'MMM d, yyyy')         -- "Jul 20, 1969"
TO_STRING(`1969-07-20T20:18Z`, 'M-d-yy')             -- "7-20-69"
TO_STRING(`1969-07-20T20:18Z`, 'MM-d-y')             -- "07-20-1969"
TO_STRING(`1969-07-20T20:18Z`, 'MMMM d, y h:m a')    -- "July 20, 1969 8:18 PM"
TO_STRING(`1969-07-20T20:18Z`, 'y-MM-dd'T' 'H:m:ssX') -- "1969-07-20T20:18:00Z"
TO_STRING(`1969-07-20T20:18+08:00Z`, 'y-MM-dd'T' 'H:m:ssX') -- "1969-07-20T20:18:00Z"
TO_STRING(`1969-07-20T20:18+08:00`, 'y-MM-dd'T' 'H:m:ssXXXX') -- "1969-07-20T20
:18:00+0800"
TO_STRING(`1969-07-20T20:18+08:00`, 'y-MM-dd'T' 'H:m:ssXXXXX') -- "1969-07-20T20
:18:00+08:00"

```

### 7.3.17 TO\_TIMESTAMP

Given a string convert it to a timestamp. This is the inverse operation of `TO_STRING`

**Signature** `TO_TIMESTAMP: String [ TimeFormatPattern ] -> Timestamp`

See definition of `TimeFormatPattern` in `TO_STRING`.

**Header** `TO_TIMESTAMP(str[ , f ])`

**Purpose** Given a string, `str`, and an optional format pattern, `f`, as a String return a timestamp whose values are extracted from `str` using `f`.

If the `<format pattern>` argument is omitted, `<string>` is assumed to be in the format of a [standard Ion timestamp](#). This is the only recommended way to parse an Ion timestamp using this function.

Zero padding is optional when using a single format symbol (e.g. `y`, `M`, `d`, `H`, `h`, `m`, `s`) but required for their zero padded variants (e.g. `yyyy`, `MM`, `dd`, `HH`, `hh`, `mm`, `ss`).

Special treatment is given to 2-digit years (format symbol `yy`). 1900 is added to values greater than or equal to 70 and 2000 is added to values less than 70.

Month names and AM/PM specifiers are case-insensitive.

## Examples

Single argument parsing an Ion timestamp:

```
TO_TIMESTAMP('2007T')           -- `2007T`
TO_TIMESTAMP('2007-02-23T12:14:33.079-08:00') -- `2007-02-23T12:14:33.079-08:00`
TO_TIMESTAMP('2016', 'y')       -- `2016T`
TO_TIMESTAMP('2016', 'yyyy')    -- `2016T`
TO_TIMESTAMP('02-2016', 'MM-yyyy') -- `2016-02T`
TO_TIMESTAMP('Feb 2016', 'MMM yyyy') -- `2016-02T`
TO_TIMESTAMP('February 2016', 'MMMM yyyy') -- `2016-02T`
```

Notes:

[All issues for PartiQL's TO\\_TIMESTAMP function.](#)

Internally, this is implemented with Java 8's `java.time` package. There are a few differences between Ion's timestamp and the `java.time` package that create a few hypothetically infrequently encountered caveats that do not really have good workarounds at this time.

- The Ion specification allows for explicitly signifying an unknown timestamp with a negative zero offset (i.e. the `-00:00` at the end of `2007-02-23T20:14:33.079-00:00`) but Java 8's `DateTimeFormatter` doesn't recognize this. **Hence, unknown offsets specified in this manner will be parsed as if they had an offset of +00:00, i.e. UTC.** To avoid this issue when parsing Ion formatted timestamps, use the single argument variant of `TO_TIMESTAMP`. There is no workaround for custom format patterns at this time.
- `DateTimeFormatter` is capable of parsing UTC offsets to the precision of seconds, but Ion Timestamp's precision for offsets is minutes. `TimestampParser` currently handles this by throwing an exception when an attempt is made to parse a timestamp with an offset that does not land on a minute boundary. For example, parsing this timestamp would throw an exception: `May 5, 2017 8:52pm +08:00:01` while `May 5, 2017 8:52pm +08:00:00` would not.
- Ion Java's Timestamp allows specification of offsets up to `+/- 23:59`, while an exception is thrown by `DateTimeFormatter` for any attempt to parse an offset greater than `+/- 18:00`. For example, attempting to parse: `May 5, 2017 8:52pm +18:01` would cause an exception to be thrown. (Note: the Ion specification does indicate minimum and maximum allowable values for offsets.) In practice this will not be an issue for systems that do not abuse the offset portion of Timestamp because real-life offsets do not exceed `+/- 12h`.

### 7.3.18 TRIM

Trims leading and/or trailing characters from a String.

**Signature** `TRIM: [ String ] String -> String`

**Header** `TRIM([[LEADING|TRAILING|BOTH r] FROM] str)`

**Purpose** Given a string, `str`, and an optional set of characters to remove, `r`, specified as a string, return the string with any character from `r` found at the beginning or end of `str` removed.

If `r` is not provided it defaults to ' '.

#### Examples

```
TRIM('      foobar      ')      -- 'foobar'
TRIM('      \tfoobar\t      ')  -- '\tfoobar\t'
TRIM(LEADING FROM '      foobar      ') -- 'foobar      '
TRIM(TRAILING FROM '      foobar      ') -- '      foobar'
TRIM(BOTH FROM '      foobar      ')  -- 'foobar'
TRIM(BOTH '☹' FROM '☹☹☹☹foobar')    -- 'foobar'
TRIM(BOTH '12' FROM '1112211foobar22211122') -- 'foobar'
```

### 7.3.19 UPPER

Given a string convert all lower case characters to upper case characters.

**Signature** `UPPER: String -> String`

**Header** `UPPER(str)`

**Purpose** Given a string, `str`, alter every upper case character is `str` to lower case. Any non-lower cases characters remain unchanged. This operation does rely on the locale specified by the runtime configuration. The implementation, currently, relies on Java's [String.toLowerCase\(\)](#) documentation.

#### Examples

```
UPPER('AbCdEfG!@#') -- 'ABCDEFG!@#'
```

### 7.3.20 UTCNOW

Returns the current time in UTC as a timestamp.

**Signature** `UTCNOW: -> Timestamp`

**Header** `UTCNOW()`

**Purpose** Return the current time in UTC as a timestamp

#### Examples



```
UTCNOW() -- 2017-10-13T16:02:11.123Z
```

### 7.3.21 UNIX\_TIMESTAMP

With no `timestamp` argument, returns the number of seconds since the last epoch ('1970-01-01 00:00:00' UTC).

With a `timestamp` argument, returns the number of seconds from the last epoch to the given `timestamp` (possibly negative).

Signature : `UNIX_TIMESTAMP: [Timestamp] -> Integer|Decimal`

Header : `UNIX_TIMESTAMP([timestamp])`

Purpose : `UNIX_TIMESTAMP()` called without a `timestamp` argument returns the number of whole seconds since the last epoch ('1970-01-01 00:00:00' UTC) as an Integer using `UTCNOW`.

`UNIX_TIMESTAMP()` called with a `timestamp` argument returns the number of seconds from the last epoch to the `timestamp` argument. If given a `timestamp` before the last epoch, `UNIX_TIMESTAMP` will return the number of seconds before the last epoch as a negative number. The return value will be a Decimal if and only if the given `timestamp` has a fractional seconds part.

Examples :

```
UNIX_TIMESTAMP()                -- 1507910531 (if current time is `2017-10-13
    T16:02:11Z`; # of seconds since last epoch as an Integer)
UNIX_TIMESTAMP(`2020T`)         -- 1577836800 (seconds from 2020 to the last
    epoch as an Integer)
UNIX_TIMESTAMP(`2020-01T`)      -- ''
UNIX_TIMESTAMP(`2020-01-01T`)   -- ''
UNIX_TIMESTAMP(`2020-01-01T00:00Z`) -- ''
UNIX_TIMESTAMP(`2020-01-01T00:00:00Z`) -- ''
UNIX_TIMESTAMP(`2020-01-01T00:00:00.0Z`) -- 1577836800. (seconds from 2020 to the last
    epoch as a Decimal)
UNIX_TIMESTAMP(`2020-01-01T00:00:00.00Z`) -- ''
UNIX_TIMESTAMP(`2020-01-01T00:00:00.000Z`) -- ''
UNIX_TIMESTAMP(`2020-01-01T00:00:00.100Z`) -- 1577836800.1
UNIX_TIMESTAMP(`1969T`)         -- -31536000 (timestamp is before last epoch)
```

### 7.3.22 FROM\_UNIXTIME

Converts the given unix epoch into a timestamp.

Signature : `FROM_UNIXTIME: Integer|Decimal -> Timestamp`

Header : `FROM_UNIXTIME(unix_timestamp)`

Purpose : When given a non-negative numeric value, returns a timestamp after the last epoch. When given a negative numeric value, returns a timestamp before the last epoch. The returned timestamp has fractional seconds depending on if the value is a decimal.

Examples :

```
FROM_UNIXTIME(-1)      -- `1969-12-31T23:59:59-00:00`      (negative unix_timestamp;  
    returns timestamp before last epoch)  
FROM_UNIXTIME(-0.1)    -- `1969-12-31T23:59:59.9-00:00`    (unix_timestamp is decimal  
    so timestamp has fractional seconds)  
FROM_UNIXTIME(0)       -- `1970-01-01T00:00:00.000-00:00`  
FROM_UNIXTIME(0.001)   -- `1970-01-01T00:00:00.001-00:00` (decimal precision to  
    fractional second precision)  
FROM_UNIXTIME(0.01)    -- `1970-01-01T00:00:00.01-00:00`  
FROM_UNIXTIME(0.1)     -- `1970-01-01T00:00:00.1-00:00`  
FROM_UNIXTIME(1)       -- `1970-01-01T00:00:01-00:00`  
FROM_UNIXTIME(1577836800) -- `2020-01-01T00:00:00-00:00`    (unix_timestamp is Integer  
    so no fractional seconds)
```

## 7.4 Math Functions

### 7.4.1 CEIL / CEILING

Returns the nearest integer greater than or equal to the input.

Signature : `CEIL/CEILING: Numeric -> Numeric`

Examples:

```
CEIL(1.1) = 2  
CEIL(-42.8) = -42  
CEIL(`+inf`) = `+inf` -- Float  
CEIL(`-inf`) = `-inf` -- Float  
CEIL(`nan`) = `nan` -- Float
```

### 7.4.2 FLOOR

Returns the nearest integer less than or equal to the input.

Signature : `FLOOR: Numeric -> Numeric`

Examples:

```
FLOOR(1.1) = 1
FLOOR(-42.8) = -43
FLOOR(`+inf`) = `+inf` -- Float
FLOOR(`-inf`) = `-inf` -- Float
FLOOR(`nan`) = `nan` -- Float
```

## References

[1] Ion Committee, “Ion Specification 1.0,” 2009. [Online]. Available: <https://amzn.github.io/ion-docs/spec.html>.